

Kubernetes is an open-source system for automating deployment, scaling, and management of containerized applications. The CNCF/Linux Foundation offers this performance-based exam which targets the developer aspect of kubernetes skills such as deploying apps, configuring apps, rolling out the application, creating persistent volumes, etc.

Since this exam is performance-based rather than just multiple choice questions just knowing the concepts are not enough, we need a lot of practice before the exam. This article helps you understand, practice and get you ready for the exam.

We are not going to discuss any concepts here, rather, I just want to create a bunch of practice questions for the CKAD exam based on the curriculum provided [here](#).

- ***Core Concepts (13%)***
- ***Multi-Container Pods (10%)***
- ***Pod Design (20%)***
- ***State Persistence (8%)***
- ***Configuration (18%)***
- ***Observability (18%)***
- ***Services and Networking (13%)***

## **Core Concepts (13%)**

Practice questions based on these concepts

- Understand Kubernetes API Primitives
- Create and Configure Basic Pods

1. ***List all the namespaces in the cluster***

```
kubectl get namespaces
```

```
kubectl get ns
```

2. ***List all the pods in all namespaces***

```
kubectl get po --all-namespaces
```

3. ***List all the pods in the particular namespace***

```
kubectl get po -n <namespace name>
```

#### ***4. List all the services in the particular namespace***

```
kubectl get svc -n <namespace name>
```

#### ***5. List all the pods showing name and namespace with a json path expression***

```
kubectl get pods  
-o=jsonpath="{.items[*]['metadata.name',  
'metadata.namespace']}"
```

#### ***6. Create an nginx pod in a default namespace and verify the pod running***

```
// creating a pod
```

```
kubect1 run nginx --image=nginx --restart=Never
```

```
// List the pod
```

```
kubect1 get po
```

## ***7. Create the same nginx pod with a yam1 file***

```
// get the yam1 file with --dry-run flag
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yam1 > nginx-pod.yam1
```

```
// cat nginx-pod.yam1
```

```
apiVersion: v1
```

kind: Pod

metadata:

creationTimestamp: null

labels:

run: nginx

name: nginx

spec:

containers:

- image: nginx

name: nginx

resources: {}

dnsPolicy: ClusterFirst

restartPolicy: Never

status: {}

```
// create a pod
```

```
kubectl create -f nginx-pod.yaml
```

***8. Output the yaml file of the pod you just created***

```
kubectl get po nginx -o yaml
```

***9. Output the yaml file of the pod you just created without the cluster-specific information***

```
kubectl get po nginx -o yaml --export
```



***10. Get the complete details of the pod you just created***

```
kubectl describe pod nginx
```

***11. Delete the pod you just created***

```
kubectl delete po nginx
```

```
kubectl delete -f nginx-pod.yaml
```

***12. Delete the pod you just created without any delay (force delete)***

```
kubectl delete po nginx --grace-period=0 --force
```

***13. Create the nginx pod with version 1.17.4  
and expose it on port 80***

```
kubect1 run nginx --image=nginx:1.17.4  
--restart=Never --port=80
```

***14. Change the Image version to 1.15-alpine  
for the pod you just created and verify the  
image version is updated***

```
kubect1 set image pod/nginx nginx=nginx:1.15-alpine
```

```
kubect1 describe po nginx
```

```
// another way it will open vi editor and change  
the version
```

```
kubectl edit po nginx
```

```
kubectl describe po nginx
```

### ***15. Change the Image version back to 1.17.1***

**for the pod you just updated and observe the changes**

```
kubectl set image pod/nginx nginx=nginx:1.17.1
```

```
kubectl describe po nginx
```

```
kubectl get po nginx -w # watch it
```

### ***16. Check the Image version without the describe command***

```
kubectl get po nginx -o  
jsonpath='{.spec.containers[].image}{"\n"}'
```

## ***17. Create the nginx pod and execute the simple shell on the pod***

```
// creating a pod
```

```
kubectl run nginx --image=nginx --restart=Never
```

```
// exec into the pod
```

```
kubectl exec -it nginx /bin/sh
```

## ***18. Get the IP Address of the pod you just created***

```
kubectl get po nginx -o wide
```

***19. Create a busybox pod and run command  
ls while creating it and check the logs***

```
kubect1 run busybox --image=busybox --restart=Never  
-- ls
```

```
kubect1 logs busybox
```

***20. If pod crashed check the previous logs of  
the pod***

```
kubect1 logs busybox -p
```

***21. Create a busybox pod with command  
sleep 3600***

```
kubect1 run busybox --image=busybox --restart=Never  
-- /bin/sh -c "sleep 3600"
```

## ***22. Check the connection of the nginx pod from the busybox pod***

```
kubectl get po nginx -o wide
```

```
// check the connection
```

```
kubectl exec -it busybox -- wget -o- <IP Address>
```

## ***23. Create a busybox pod and echo message ‘How are you’ and delete it manually***

```
kubectl run busybox --image=nginx --restart=Never  
-it -- echo "How are you"
```

```
kubectl delete po busybox
```

## ***24. Create a busybox pod and echo message***

***‘How are you’ and have it deleted***

***immediately***

```
// notice the --rm flag
```

```
kubect1 run busybox --image=nginx --restart=Never  
-it --rm -- echo "How are you"
```

## ***25. Create an nginx pod and list the pod with***

***different levels of verbosity***

```
// create a pod
```

```
kubect1 run nginx --image=nginx --restart=Never  
--port=80
```

```
// List the pod with different verbosity
```

```
kubectl get po nginx --v=7
```

```
kubectl get po nginx --v=8
```

```
kubectl get po nginx --v=9
```

## ***26. List the nginx pod with custom columns***

### ***POD\_NAME and POD\_STATUS***

```
kubectl get po  
-o=custom-columns="POD_NAME:.metadata.name,  
POD_STATUS:.status.containerStatuses[].state"
```

## ***27. List all the pods sorted by name***

```
kubectl get pods --sort-by=.metadata.name
```



***28. List all the pods sorted by created timestamp***

```
kubectl get  
pods--sort-by=.metadata.creationTimestamp
```

## **Multi-Container Pods (10%)**

Practice questions based on these concepts

- Understand multi-container pod design patterns (eg: ambassador, adaptor, sidecar)

***29. Create a Pod with three busy box containers with commands “ls; sleep 3600;”, “echo Hello World; sleep 3600;” and***

***“echo this is the third container; sleep 3600”***

***respectively and check the status***

```
// first create single container pod with dry run  
flag
```

```
kubect1 run busybox --image=busybox --restart=Never  
--dry-run -o yaml -- bin/sh -c "sleep 3600; ls" >  
multi-container.yaml
```

```
// edit the pod to following yaml and create it
```

```
kubect1 create -f multi-container.yaml
```

```
kubect1 get po busybox  
multi-container pod
```

***30. Check the logs of each container that you just created***

```
kubectl logs busybox -c busybox1
```

```
kubectl logs busybox -c busybox2
```

```
kubectl logs busybox -c busybox3
```

***31. Check the previous logs of the second container busybox2 if any***

```
kubectl logs busybox -c busybox2 --previous
```

***32. Run command ls in the third container busybox3 of the above pod***

```
kubectl exec busybox -c busybox3 -- ls
```

***33. Show metrics of the above pod  
containers and puts them into the file.log  
and verify***

```
kubect1 top pod busybox --containers
```

```
// putting them into file
```

```
kubect1 top pod busybox --containers > file.log
```

```
cat file.log
```

***34. Create a Pod with main container  
busybox and which executes this “while  
true; do echo ‘Hi I am from Main container’***

**>> /var/log/index.html; sleep 5; done” and  
with sidecar container with nginx image  
which exposes on port 80. Use emptyDir  
Volume and mount this volume on path  
/var/log for busybox and on path  
/usr/share/nginx/html for nginx container.  
Verify both containers are running.**

```
// create an initial yaml file with this
```

```
kubectl run multi-cont-pod --image=busybox  
--restart=Never --dry-run -o yaml >  
multi-container.yaml
```

```
// edit the yml as below and create it
```

```
kubectl create -f multi-container.yaml
```

```
kubectl get po multi-cont-pod  
multi-container.yaml
```

***35. Exec into both containers and verify that  
main.txt exist and query the main.txt from  
sidecar container with `curl localhost`***

```
// exec into main container
```

```
kubectl exec -it multi-cont-pod -c main-container  
-- sh
```

```
cat /var/log/main.txt
```

```
// exec into sidecar container
```

```
kubectl exec -it multi-cont-pod -c  
sidecar-container -- sh
```

```
cat /usr/share/nginx/html/index.html
```

```
// install curl and get default page
```

```
kubectl exec -it multi-cont-pod -c  
sidecar-container -- sh
```

```
# apt-get update && apt-get install -y curl
```

```
# curl localhost
```

---

## Pod Design (20%)

Practice questions based on these concepts

- Understand how to use Labels, Selectors and Annotations
- Understand Deployments and how to perform rolling updates
- Understand Deployments and how to perform rollbacks
- Understand Jobs and CronJobs

### ***36. Get the pods with label information***

```
kubectl get pods --show-labels
```

### ***37. Create 5 nginx pods in which two of them is labeled env=prod and three of them is labeled env=dev***



```
kubect1 run nginx-dev1 --image=nginx  
--restart=Never --labels=env=dev
```

```
kubect1 run nginx-dev2 --image=nginx  
--restart=Never --labels=env=dev
```

```
kubect1 run nginx-dev3 --image=nginx  
--restart=Never --labels=env=dev
```

```
kubect1 run nginx-prod1 --image=nginx  
--restart=Never --labels=env=prod
```

```
kubect1 run nginx-prod2 --image=nginx  
--restart=Never --labels=env=prod
```

***38. Verify all the pods are created with  
correct labels***

```
kubect1 get pods --show-labels
```

***39. Get the pods with label env=dev***

```
kubectl get pods -l env=dev
```

***40. Get the pods with label env=dev and also  
output the labels***

```
kubectl get pods -l env=dev --show-labels
```

***41. Get the pods with label env=prod***

```
kubectl get pods -l env=prod
```

***42. Get the pods with label env=prod and  
also output the labels***

```
kubectl get pods -l env=prod --show-labels
```

***43. Get the pods with label env***

```
kubectl get pods -L env
```

***44. Get the pods with labels env=dev and env=prod***

```
kubectl get pods -l 'env in (dev,prod)'
```

***45. Get the pods with labels env=dev and env=prod and output the labels as well***

```
kubectl get pods -l 'env in (dev,prod)'  
--show-labels
```

***46. Change the label for one of the pod to env=uat and list all the pods to verify***

```
kubectl label pod/nginx-dev3 env=uat --overwrite
```

```
kubectl get pods --show-labels
```

***47. Remove the labels for the pods that we created now and verify all the labels are removed***

```
kubectl label pod nginx-dev{1..3} env-
```

```
kubectl label pod nginx-prod{1..2} env-
```

```
kubectl get po --show-labels
```

***48. Let's add the label app=nginx for all the pods and verify***

```
kubectl label pod nginx-dev{1..3} app=nginx
```

```
kubectl label pod nginx-prod{1..2} app=nginx
```

```
kubectl get po --show-labels
```

***49. Get all the nodes with labels (if using minikube you would get only master node)***

```
kubectl get nodes --show-labels
```

***50. Label the node (minikube if you are using) nodeName=nginxnode***

```
kubectl label node minikube nodeName=nginxnode
```

***51. Create a Pod that will be deployed on this node with the label nodeName=nginxnode***

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > pod.yaml
```

```
// add the nodeSelector like below and create the  
pod
```

```
kubect1 create -f pod.yaml  
pod.yaml
```

## ***52. Verify the pod that it is scheduled with the node selector***

```
kubect1 describe po nginx | grep Node-Selectors
```

## ***53. Verify the pod nginx that we just created has this label***

```
kubect1 describe po nginx | grep Labels
```

## ***54. Annotate the pods with name=webapp***

```
kubectl annotate pod nginx-dev{1..3} name=webapp
```

```
kubectl annotate pod nginx-prod{1..2} name=webapp
```

## ***55. Verify the pods that have been annotated correctly***

```
kubectl describe po nginx-dev{1..3} | grep -i  
annotations
```

```
kubectl describe po nginx-prod{1..2} | grep -i  
annotations
```

## ***56. Remove the annotations on the pods and verify***

```
kubectl annotate pod nginx-dev{1..3} name-
```

```
kubectl annotate pod nginx-prod{1..2} name-
```

```
kubectl describe po nginx-dev{1..3} | grep -i  
annotations
```

```
kubectl describe po nginx-prod{1..2} | grep -i  
annotations
```

***57. Remove all the pods that we created so far***

```
kubectl delete po --all
```

***58. Create a deployment called webapp with image nginx with 5 replicas***

```
kubectl create deploy webapp --image=nginx  
--dry-run -o yaml > webapp.yaml
```



```
// change the replicas to 5 in the yaml and create  
it
```

```
kubect1 create -f webapp.yaml  
webapp.yaml
```

### ***59. Get the deployment you just created with labels***

```
kubect1 get deploy webapp --show-labels
```

### ***60. Output the yaml file of the deployment you just created***

```
kubect1 get deploy webapp -o yaml
```

### ***61. Get the pods of this deployment***

```
// get the label of the deployment
```

```
kubectl get deploy --show-labels
```

```
// get the pods with that label
```

```
kubectl get pods -l app=webapp
```

## ***62. Scale the deployment from 5 replicas to 20 replicas and verify***

```
kubectl scale deploy webapp --replicas=20
```

```
kubectl get po -l app=webapp
```

## ***63. Get the deployment rollout status***

```
kubectl rollout status deploy webapp
```

***64. Get the replicaset that created with this deployment***

```
kubectl get rs -l app=webapp
```

***65. Get the yaml of the replicaset and pods of this deployment***

```
kubectl get rs -l app=webapp -o yaml
```

```
kubectl get po -l app=webapp -o yaml
```

***66. Delete the deployment you just created and watch all the pods are also being deleted***

```
kubectl delete deploy webapp
```

```
kubectl get po -l app=webapp -w
```

***67. Create a deployment of webapp with  
image nginx:1.17.1 with container port 80  
and verify the image version***

```
kubectl create deploy webapp --image=nginx:1.17.1  
--dry-run -o yaml > webapp.yaml
```

```
// add the port section and create the deployment
```

```
kubectl create -f webapp.yaml
```

```
// verify
```

```
kubectl describe deploy webapp | grep Image  
webapp.yaml
```

## ***68. Update the deployment with the image version 1.17.4 and verify***

```
kubectl set image deploy/webapp nginx=nginx:1.17.4
```

```
kubectl describe deploy webapp | grep Image
```

## ***69. Check the rollout history and make sure everything is ok after the update***

```
kubectl rollout history deploy webapp
```

```
kubectl get deploy webapp --show-labels
```

```
kubectl get rs -l app=webapp
```

```
kubectl get po -l app=webapp
```

***70. Undo the deployment to the previous version 1.17.1 and verify Image has the previous version***

```
kubectl rollout undo deploy webapp
```

```
kubectl describe deploy webapp | grep Image
```

***71. Update the deployment with the image version 1.16.1 and verify the image and also check the rollout history***

```
kubectl set image deploy/webapp nginx=nginx:1.16.1
```

```
kubectl describe deploy webapp | grep Image
```

```
kubectl rollout history deploy webapp
```

## ***72. Update the deployment to the Image***

### ***1.17.1 and verify everything is ok***

```
kubectl rollout undo deploy webapp --to-revision=3
```

```
kubectl describe deploy webapp | grep Image
```

```
kubectl rollout status deploy webapp
```

***73. Update the deployment with the wrong image version 1.100 and verify something is wrong with the deployment***

```
kubectl set image deploy/webapp nginx=nginx:1.100
```

```
kubectl rollout status deploy webapp (still pending state)
```

```
kubectl get pods (ImagePullErr)
```

***74. Undo the deployment with the previous version and verify everything is Ok***

```
kubectl rollout undo deploy webapp
```

```
kubectl rollout status deploy webapp
```



```
kubectl get pods
```

***75. Check the history of the specific revision  
of that deployment***

```
kubectl rollout history deploy webapp --revision=7
```

***76. Pause the rollout of the deployment***

```
kubectl rollout pause deploy webapp
```

***77. Update the deployment with the image  
version latest and check the history and  
verify nothing is going on***

```
kubectl set image deploy/webapp nginx=nginx:latest
```

```
kubectl rollout history deploy webapp (No new revision)
```

## ***78. Resume the rollout of the deployment***

```
kubectl rollout resume deploy webapp
```

## ***79. Check the rollout history and verify it has the new version***

```
kubectl rollout history deploy webapp
```

```
kubectl rollout history deploy webapp --revision=9
```

## ***80. Apply the autoscaling to this deployment with minimum 10 and maximum 20 replicas and target CPU of***

***85% and verify hpa is created and replicas  
are increased to 10 from 1***

```
kubectl autoscale deploy webapp --min=10 --max=20  
--cpu-percent=85
```

```
kubectl get hpa
```

```
kubectl get pod -l app=webapp
```

***81. Clean the cluster by deleting deployment  
and hpa you just created***

```
kubectl delete deploy webapp
```

```
kubectl delete hpa webapp
```

***82. Create a Job with an image node which prints node version and also verifies there is a pod created for this job***

```
kubect1 create job nodeversion --image=node -- node  
-v
```

```
kubect1 get job -w
```

```
kubect1 get pod
```

***83. Get the logs of the job just created***

```
kubect1 logs <pod name> // created from the job
```

***84. Output the yaml file for the Job with the image busybox which echos “Hello I am from job”***

```
kubect1 create job hello-job --image=busybox  
--dry-run -o yaml -- echo "Hello I am from job"
```

***85. Copy the above YAML file to hello-job.yaml file and create the job***

```
kubect1 create job hello-job --image=busybox  
--dry-run -o yaml -- echo "Hello I am from job" >  
hello-job.yaml
```

```
kubect1 create -f hello-job.yaml
```

***86. Verify the job and the associated pod is created and check the logs as well***

```
kubectl get job
```

```
kubectl get po
```

```
kubectl logs hello-job-*
```

## ***87. Delete the job we just created***

```
kubectl delete job hello-job
```

## ***88. Create the same job and make it run 10 times one after one***

```
kubectl create job hello-job --image=busybox  
--dry-run -o yaml -- echo "Hello I am from job" >  
hello-job.yaml
```

```
// edit the yaml file to add completions: 10
```

```
kubectl create -f hello-job.yaml  
hello-job.yaml
```

***89. Watch the job that runs 10 times one by one and verify 10 pods are created and delete those after it's completed***

```
kubectl get job -w
```

```
kubectl get po
```

```
kubectl delete job hello-job
```

***90. Create the same job and make it run 10 times parallel***

```
kubectl create job hello-job --image=busybox
--dry-run -o yaml -- echo "Hello I am from job" >
hello-job.yaml
```

```
// edit the yaml file to add parallelism: 10
```

```
kubectl create -f hello-job.yaml
hello-job.yaml
```

***91. Watch the job that runs 10 times***

***parallelly and verify 10 pods are created***

***and delete those after it's completed***

```
kubectl get job -w
```

```
kubectl get po
```

```
kubectl delete job hello-job
```



***92. Create a Cronjob with busybox image  
that prints date and hello from kubernetes  
cluster message for every minute***

```
kubect1 create cronjob date-job --image=busybox  
--schedule="*/1 * * * *" -- bin/sh -c "date; echo  
Hello from kubernetes cluster"
```

***93. Output the YAML file of the above  
cronjob***

```
kubect1 get cj date-job -o yaml
```

***94. Verify that CronJob creating a separate  
job and pods for every minute to run and  
verify the logs of the pod***

```
kubect1 get job
```

```
kubectl get po
```

```
kubectl logs date-job-<jobid>--<pod>
```

***95. Delete the CronJob and verify all the associated jobs and pods are also deleted.***

```
kubectl delete cj date-job
```

```
// verify pods and jobs
```

```
kubectl get po
```

```
kubectl get job
```

---

## **State Persistence (8%)**

Practice questions based on these concepts

- Understand PersistentVolumeClaims for  
Storage

***96. List Persistent Volumes in the cluster***

```
kubectl get pv
```

***97. Create a hostPath PersistentVolume***

***named task-pv-volume with storage 10Gi,***

***access modes ReadWriteOnce,***

***storageClassName manual, and volume at***

***/mnt/data and verify***

```
kubectl create -f task-pv-volume.yaml
```

```
kubect1 get pv  
task-pv-volume.yaml
```

***98. Create a PersistentVolumeClaim of at  
least 3Gi storage and access mode  
ReadWriteOnce and verify status is Bound***

```
kubect1 create -f task-pv-claim.yaml
```

```
kubect1 get pvc  
task-pv-claim.yaml
```

***99. Delete persistent volume and  
PersistentVolumeClaim we just created***

```
kubect1 delete pvc task-pv-claim
```

```
kubect1 delete pv task-pv-volume
```

***100. Create a Pod with an image Redis and configure a volume that lasts for the lifetime of the Pod***

```
// emptyDir is the volume that lasts for the life  
of the pod
```

```
kubectl create -f redis-storage.yaml
```

***101. Exec into the above pod and create a file named file.txt with the text 'This is called the file' in the path /data/redis and open another tab and exec again with the same pod and verifies file exist in the same path.***

```
// first terminal
```

```
kubectl exec -it redis-storage /bin/sh
```

```
cd /data/redis
```

```
echo 'This is called the file' > file.txt
```

```
//open another tab
```

```
kubectl exec -it redis-storage /bin/sh
```

```
cat /data/redis/file.txt
```

***102. Delete the above pod and create again  
from the same yaml file and verifies there is  
no file.txt in the path /data/redis***

```
kubectl delete pod redis
```

```
kubectl create -f redis-storage.yaml
```

```
kubectl exec -it redis-storage /bin/sh
```

```
cat /data/redis/file.txt // file doesn't exist
```

***103. Create PersistentVolume named task-pv-volume with storage 10Gi, access modes ReadWriteOnce, storageClassName manual, and volume at /mnt/data and Create a PersistentVolumeClaim of at least 3Gi storage and access mode ReadWriteOnce and verify status is Bound***

```
kubectl create -f task-pv-volume.yaml
```

```
kubectl create -f task-pv-claim.yaml
```

```
kubectl get pv
```

```
kubectl get pvc
```

## ***104. Create an nginx pod with containerPort 80 and with a PersistentVolumeClaim***

***task-pv-claim and has a mount path***

***"/usr/share/nginx/html"***

```
kubectl create -f task-pv-pod.yaml  
task-pv-pod.yaml
```

---

## **Configuration (18%)**

Practice questions based on these concepts



- Understand ConfigMaps
- Understand SecurityContexts
- Define an application's resource requirements
- Create & Consume Secrets
- Understand ServiceAccounts

***105. List all the configmaps in the cluster***

```
kubectl get cm
```

or

```
kubectl get configmap
```

***106. Create a configmap called  
myconfigmap with literal value  
appname=myapp***

```
kubectl create cm myconfigmap  
--from-literal=appname=myapp
```

***107. Verify the configmap we just created  
has this data***

```
// you will see under data
```

```
kubectl get cm -o yaml
```

or

```
kubectl describe cm
```

***108. delete the configmap myconfigmap we just created***

```
kubectl delete cm myconfigmap
```

***109. Create a file called config.txt with two values key1=value1 and key2=value2 and verify the file***

```
cat >> config.txt << EOF
```

```
key1=value1
```

```
key2=value2
```

```
EOF
```

```
cat config.txt
```

***110. Create a configmap named keyvalcfgmap and read data from the file config.txt and verify that configmap is created correctly***

```
kubectl create cm keyvalcfgmap  
--from-file=config.txt
```

```
kubectl get cm keyvalcfgmap -o yaml
```

***111. Create an nginx pod and load environment values from the above configmap keyvalcfgmap and exec into the***

## ***pod and verify the environment variables and delete the pod***

```
// first run this command to save the pod yml
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx-pod.yml
```

```
// edit the yml to below file and create
```

```
kubect1 create -f nginx-pod.yml
```

```
// verify
```

```
kubect1 exec -it nginx -- env
```

```
kubect1 delete po nginx  
nginx-pod.yml
```

***112. Create an env file file.env with  
var1=val1 and create a configmap  
envcfgmap from this env file and verify the  
configmap***

```
echo var1=val1 > file.env
```

```
cat file.env
```

```
kubectl create cm envcfgmap  
--from-env-file=file.env
```

```
kubectl get cm envcfgmap -o yaml --export
```

***113. Create an nginx pod and load  
environment values from the above***

***configmap envcfgmap and exec into the pod  
and verify the environment variables and  
delete the pod***

```
// first run this command to save the pod yml
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx-pod.yml
```

```
// edit the yml to below file and create
```

```
kubect1 create -f nginx-pod.yml
```

```
// verify
```

```
kubect1 exec -it nginx -- env
```

```
kubect1 delete po nginx  
nginx-pod.yaml
```

***114. Create a configmap called cfgvolume  
with values var1=val1, var2=val2 and create  
an nginx pod with volume nginx-volume  
which reads data from this configmap  
cfgvolume and put it on the path /etc/cfg***

```
// first create a configmap cfgvolume
```

```
kubect1 create cm cfgvolume  
--from-literal=var1=val1 --from-literal=var2=val2
```

```
// verify the configmap
```

```
kubect1 describe cm cfgvolume
```



```
// create the config map
```

```
kubectl create -f nginx-volume.yml
```

```
// exec into the pod
```

```
kubectl exec -it nginx -- /bin/sh
```

```
// check the path
```

```
cd /etc/cfg
```

```
ls  
nginx-volume.yml
```

***115. Create a pod called secbusybox with the image busybox which executes command sleep 3600 and makes sure any Containers in the Pod, all processes run with user ID 1000 and with group id 2000 and verify.***

```
// create yaml file with dry-run
```

```
kubect1 run secbusybox --image=busybox  
--restart=Never --dry-run -o yaml -- /bin/sh -c  
"sleep 3600;" > busybox.yaml
```

```
// edit the pod like below and create
```

```
kubect1 create -f busybox.yaml
```

```
// verify
```

```
kubectll exec -it secbusybox -- sh
```

```
id // it will show the id and group  
busybox.yml
```

***116. Create the same pod as above this time  
set the securityContext for the container as  
well and verify that the securityContext of  
container overrides the Pod level  
securityContext.***

```
// create yaml file with dry-run
```

```
kubectll run secbusybox --image=busybox  
--restart=Never --dry-run -o yaml -- /bin/sh -c  
"sleep 3600;" > busybox.yml
```

```
// edit the pod like below and create
```

```
kubectl create -f busybox.yml
```

```
// verify
```

```
kubectl exec -it secbusybox -- sh
```

```
id // you can see container securityContext  
overrides the Pod level  
busybox.yml
```

***117. Create pod with an nginx image and***

***configure the pod with capabilities `NET_ADMIN`***

***and `SYS_TIME` verify the capabilities***

```
// create the yaml file
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yaml
```

```
// edit as below and create pod
```

```
kubect1 create -f nginx.yaml
```

```
// exec and verify
```

```
kubect1 exec -it nginx -- sh
```

```
cd /proc/1
```

```
cat status
```

```
// you should see these values
```

```
CapPrm: 00000000aa0435fb
```

```
CapEff: 00000000aa0435fb  
nginx.yml
```

***118. Create a Pod nginx and specify a memory request and a memory limit of 100Mi and 200Mi respectively.***

```
// create a yml file
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yml
```

```
// add the resources section and create
```

```
kubect1 create -f nginx.yml
```

```
// verify
```

```
kubect1 top pod  
nginx.yml
```

***119. Create a Pod nginx and specify a CPU request and a CPU limit of 0.5 and 1 respectively.***

```
// create a yml file
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yml
```

```
// add the resources section and create
```

```
kubectl create -f nginx.yml
```

```
// verify
```

```
kubectl top pod  
nginx.yml
```

***120. Create a Pod nginx and specify both CPU, memory requests and limits together and verify.***

```
// create a yml file
```

```
kubectl run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yml
```

```
// add the resources section and create
```



```
kubectl create -f nginx.yml
```

```
// verify
```

```
kubectl top pod  
nginx.yml
```

***121. Create a Pod nginx and specify a memory request and a memory limit of 100Gi and 200Gi respectively which is too big for the nodes and verify pod fails to start because of insufficient memory***

```
// create a yml file
```

```
kubectl run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yml
```

```
// add the resources section and create
```

```
kubectl create -f nginx.yml
```

```
// verify
```

```
kubectl describe po nginx // you can see pending  
state  
nginx.yml
```

## ***122. Create a secret mysecret with values***

***user=myuser and password=mypassword***

```
kubectl create secret generic my-secret  
--from-literal=username=user  
--from-literal=password=mypassword
```

## ***123. List the secrets in all namespaces***

```
kubect1 get secret --all-namespaces
```

***124. Output the yaml of the secret created above***

```
kubect1 get secret my-secret -o yaml
```

***125. Create an nginx pod which reads username as the environment variable***

```
// create a yaml file
```

```
kubect1 run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yaml
```

```
// add env section below and create
```

```
kubect1 create -f nginx.yaml
```

```
//verify
```

```
kubectl exec -it nginx -- env  
nginx.yml
```

## ***126. Create an nginx pod which loads the secret as environment variables***

```
// create a yml file
```

```
kubectl run nginx --image=nginx --restart=Never  
--dry-run -o yaml > nginx.yml
```

```
// add env section below and create
```

```
kubectl create -f nginx.yml
```

```
//verify
```

```
kubectl exec -it nginx -- env  
nginx.yml
```

***127. List all the service accounts in the  
default namespace***

```
kubectl get sa
```

***128. List all the service accounts in all  
namespaces***

```
kubectl get sa --all-namespaces
```

***129. Create a service account called admin***

```
kubectl create sa admin
```

***130. Output the YAML file for the service account we just created***

```
kubectl get sa admin -o yaml
```

***131. Create a busybox pod which executes this command sleep 3600 with the service account admin and verify***

```
kubectl run busybox --image=busybox --restart=Never  
--dry-run -o yaml -- /bin/sh -c "sleep 3600" >  
busybox.yaml
```

```
kubectl create -f busybox.yaml
```

```
// verify
```

```
kubectl describe po busybox  
busybox.yml
```

---

## Observability (18%)

Practice questions based on these concepts

- Understand LivenessProbes and ReadinessProbes
- Understand Container Logging
- Understand how to monitor applications in kubernetes
- Understand Debugging in Kubernetes

***132. Create an nginx pod with containerPort***

***80 and it should only receive traffic only it***

***checks the endpoint / on port 80 and verify  
and delete the pod.***

```
kubect1 run nginx --image=nginx --restart=Never  
--port=80 --dry-run -o yaml > nginx-pod.yaml
```

```
// add the readinessProbe section and create
```

```
kubect1 create -f nginx-pod.yaml
```

```
// verify
```

```
kubect1 describe pod nginx | grep -i readiness
```

```
kubect1 delete po nginx  
nginx-pod.yaml
```



***133. Create an nginx pod with containerPort 80 and it should check the pod running at endpoint / healthz on port 80 and verify and delete the pod.***

```
kubectrl run nginx --image=nginx --restart=Never  
--port=80 --dry-run -o yaml > nginx-pod.yaml
```

```
// add the livenessProbe section and create
```

```
kubectrl create -f nginx-pod.yaml
```

```
// verify
```

```
kubectrl describe pod nginx | grep -i readiness
```

```
kubect1 delete po nginx  
nginx-pod.yaml
```

***134. Create an nginx pod with containerPort 80 and it should check the pod running at endpoint /healthz on port 80 and it should only receive traffic only it checks the endpoint / on port 80. verify the pod.***

```
kubect1 run nginx --image=nginx --restart=Never  
--port=80 --dry-run -o yaml > nginx-pod.yaml
```

```
// add the livenessProbe and readiness section and  
create
```

```
kubect1 create -f nginx-pod.yaml
```

```
// verify
```

```
kubectl describe pod nginx | grep -i readiness
```

```
kubectl describe pod nginx | grep -i liveness  
nginx-pod.yaml
```

***135. Check what all are the options that we  
can configure with readiness and liveness  
probes***

```
kubectl explain Pod.spec.containers.livenessProbe
```

```
kubectl explain Pod.spec.containers.readinessProbe
```

***136. Create the pod nginx with the above liveness and readiness probes so that it should wait for 20 seconds before it checks liveness and readiness probes and it should check every 25 seconds.***

```
kubect1 create -f nginx-pod.yaml  
nginx-pod.yaml
```

***137. Create a busybox pod with this command “echo I am from busybox pod; sleep 3600;” and verify the logs.***

```
kubect1 run busybox --image=busybox --restart=Never  
-- /bin/sh -c "echo I am from busybox pod; sleep  
3600;"
```

```
kubectl logs busybox
```

***138. copy the logs of the above pod to the  
busybox-logs.txt and verify***

```
kubectl logs busybox > busybox-logs.txt
```

```
cat busybox-logs.txt
```

***139. List all the events sorted by timestamp  
and put them into file.log and verify***

```
kubectl get events  
--sort-by=.metadata.creationTimestamp
```

```
// putting them into file.log
```

```
kubect1 get events  
--sort-by=.metadata.creationTimestamp > file.log
```

```
cat file.log
```

***140. Create a pod with an image alpine  
which executes this command "while true;  
do echo 'Hi I am from alpine'; sleep 5; done"  
and verify and follow the logs of the pod.***

```
// create the pod
```

```
kubect1 run hello --image=alpine --restart=Never  
-- /bin/sh -c "while true; do echo 'Hi I am from  
Alpine'; sleep 5;done"
```

```
// verify and follow the logs
```

```
kubectl logs --follow hello
```

**141. Create the pod with this `kubectl create -f`**  
**`https://gist.githubusercontent.com/bbachi/212168375b39e36e2e2984c097167b00/raw/1fd63509c3ae3a3d3da844640fb4cca744543c1c/not-running.yml`****. The pod is not in the**  
**running state. Debug it.**

```
// create the pod
```

```
kubectl create -f  
https://gist.githubusercontent.com/bbachi/212168375b39e36e2e2984c097167b00/raw/1fd63509c3ae3a3d3da844640fb4cca744543c1c/not-running.yml
```

```
// get the pod
```

```
kubect1 get pod not-running
```

```
kubect1 describe po not-running
```

```
// it clearly says ImagePullBackOff something wrong  
with image
```

```
kubect1 edit pod not-running // it will open vim  
editor
```

or

```
kubect1 set image pod/not-running not-running=nginx
```

***142. This following yaml creates 4  
namespaces and 4 pods. One of the pod in***



***one of the namespaces are not in the  
running state. Debug and fix it.***

***[https://gist.githubusercontent.com/bbachi/  
1f001f10337234d46806929d12245397/raw/8  
4b7295fb077f15de979fec5b3f7a13fc69c6d83/  
problem-pod.yaml](https://gist.githubusercontent.com/bbachi/1f001f10337234d46806929d12245397/raw/84b7295fb077f15de979fec5b3f7a13fc69c6d83/problem-pod.yaml).***

```
kubect1 create -f  
https://gist.githubusercontent.com/bbachi/1f001f103  
37234d46806929d12245397/raw/84b7295fb077f15de979fec  
5b3f7a13fc69c6d83/problem-pod.yaml
```

```
// get all the pods in all namespaces
```

```
kubect1 get po --all-namespaces
```

```
// find out which pod is not running
```

```
kubect1 get po -n namespace2
```

```
// update the image
```

```
kubect1 set image pod/pod2 pod2=nginx -n namespace2
```

```
// verify again
```

```
kubect1 get po -n namespace2
```

***143. Get the memory and CPU usage of all the pods and find out top 3 pods which have the highest usage and put them into the cpu-usage.txt file***

```
// get the top 3 hungry pods
```

```
kubectl top pod --all-namespaces | sort --reverse  
--key 3 --numeric | head -3
```

```
// putting into file
```

```
kubectl top pod --all-namespaces | sort --reverse  
--key 3 --numeric | head -3 > cpu-usage.txt
```

```
// verify
```

```
cat cpu-usage.txt
```

---

## **Services and Networking (13%)**

Practice questions based on these concepts

- Understand Services
- Demonstrate a basic understanding of NetworkPolicies

***144. Create an nginx pod with a yaml file  
with label my-nginx and expose the port 80***

```
kubectl run nginx --image=nginx --restart=Never  
--port=80 --dry-run -o yaml > nginx.yaml
```

```
// edit the label app: my-nginx and create the pod
```

```
kubectl create -f nginx.yaml  
nginx.yaml
```

***145. Create the service for this nginx pod  
with the pod selector app: my-nginx***

```
// create the below service
```

```
kubectl create -f nginx-svc.yaml  
nginx-svc.yaml
```

***146. Find out the label of the pod and verify  
the service has the same label***

```
// get the pod with labels
```

```
kubectl get po nginx --show-labels
```

```
// get the service and chekc the selector column
```

```
kubectl get svc my-service -o wide
```

***147. Delete the service and create the service with kubectl expose command and verify the label***

```
// delete the service
```

```
kubectl delete svc my-service
```

```
// create the service again
```

```
kubectl expose po nginx --port=80  
--target-port=9376
```

```
// verify the label
```

```
kubectl get svc -l app=my-nginx
```

***148. Delete the service and create the service again with type NodePort***

```
// delete the service
```

```
kubectl delete svc nginx
```

```
// create service with expose command
```

```
kubectl expose po nginx --port=80 --type=NodePort
```

***149. Create the temporary busybox pod and hit the service. Verify the service that it should return the nginx page index.html.***

```
// get the clusterIP from this command
```

```
kubectl get svc nginx -o wide
```

```
// create temporary busybox to check the nodeport
```

```
kubectl run busybox --image=busybox --restart=Never  
-it --rm -- wget -o- <Cluster IP>:80
```

## ***150. Create a NetworkPolicy which denies all ingress traffic***

```
apiVersion: networking.k8s.io/v1
```

```
kind: NetworkPolicy
```

```
metadata:
```

```
  name: default-deny
```

```
spec:
```



```
podSelector: {}
```

```
policyTypes:
```

```
- Ingress
```

---

## Conclusion

CKAD is a performance-based exam and it's all about completing 19 questions within 2 hours. We need a lot of practice for it. These 150 questions give you enough practice for the exam. The more you practice the more comfortable